

Comprehensive Technical Report on Informix Replication Strategies

Date: 2026-05-12

Report Objective: This document provides a comprehensive technical analysis of strategies for replicating large volumes of data (approximately 15TB) from one Informix database server to another, with a specific focus on achieving real-time or continuous replication while minimizing downtime.

Executive Summary

Replicating a 15-terabyte Informix database in real-time presents significant technical challenges related to performance, data consistency, and initial setup. This report analyzes the native replication technologies provided by IBM Informix—Enterprise Replication (ER), High-Availability Data Replication (HDR), and Remote Standalone Secondary (RSS)—as well as Change Data Capture (CDC) mechanisms and leading third-party solutions.

For the specified requirements of continuous, large-volume, Informix-to-Informix replication with high data consistency, a hybrid approach is recommended. This approach combines **High-Availability Data Replication (HDR) in asynchronous mode** for a robust, physically identical failover copy with **Enterprise Replication (ER)** for granular, table-level replication to support active query workloads on the target. The initial 15TB data synchronization should be performed using the Informix **High-Performance Loader (HPL)** in conjunction with a full system backup and restore to minimize downtime.

This dual strategy leverages HDR's efficiency in block-level replication for disaster recovery and ER's flexibility for creating an active, readable secondary server. This combination provides a resilient, high-performance, and functionally rich replication environment capable of handling the demands of a 15TB dataset. Careful planning of network infrastructure, logical log capacity, and disk I/O is critical to the success of this strategy.

1. Overview of Informix Replication Technologies

IBM Informix offers a suite of powerful and mature replication technologies, each designed to address specific business needs, from high availability and disaster recovery to data distribution and workload balancing. Understanding the fundamental architecture of each is crucial for selecting the appropriate strategy for a large-scale, real-time replication project. The primary native methods include Enterprise Replication (ER), which provides asynchronous, transaction-based replication at the logical level; High-Availability Data Replication (HDR), which offers physical log-based replication for creating a hot standby; and Remote Standalone Secondary (RSS), an extension of HDR for asynchronous disaster recovery scenarios. Additionally, Informix provides Change Data Capture (CDC) and streaming mechanisms, like Smart Triggers, for pushing targeted data changes to external applications. Each of these technologies presents a different balance of performance, data consistency, system overhead, and administrative complexity.

2. Enterprise Replication (ER)

Enterprise Replication (ER) is a cornerstone of Informix's data distribution capabilities, operating as an asynchronous, log-based replication solution. It is designed to move data between geographically dispersed Informix servers, supporting complex topologies like update-anywhere, primary-target, and hierarchical trees. This flexibility makes it a powerful tool for data consolidation, distributing workloads, and propagating information across an enterprise.

Architecture and Capabilities

The architecture of ER is fundamentally rooted in capturing transactions from the logical log of a source server. When a transaction is committed on a table defined for replication, ER captures the relevant log records, stores them in a reliable send queue, and transmits them to one or more target servers. Upon arrival, the transactions are placed in a receive queue and then applied to the corresponding target tables. This entire process is asynchronous, meaning the commit on the source server does not wait for the transaction to be applied on the targets, which minimizes performance impact on the primary OLTP system.

ER is highly configurable, allowing for both full-table replication and partial replication of rows and columns using fragmentation criteria or `WHERE` clauses. It supports update-anywhere configurations where multiple servers can accept updates to the same data, with built-in conflict detection and resolution mechanisms. These mechanisms can be configured based on timestamps or user-defined procedures, ensuring data integrity in distributed environments. This architectural flexibility allows ER to be used for more than just disaster recovery; it is well-suited for creating active-active clusters or off-loading reporting and business intelligence queries to a secondary server that is continuously updated.

Performance Considerations for 15TB

Replicating a 15TB database with Enterprise Replication introduces significant performance considerations. The volume of transactional data generated by such a large database will place immense strain on the ER infrastructure. ER's performance is inherently tied to transaction throughput, logical log velocity, network bandwidth, and the I/O capacity of the disks hosting the send and receive queues. For a 15TB system, the logical logs will fill rapidly, and ER must be able to process these logs efficiently to avoid falling behind. Adequate disk space for the spooling areas (send/receive queues) is critical to prevent replication from halting due to backlogs, especially during network interruptions or periods of high transactional activity. IBM Informix 15 brought significant scalability enhancements, including relaxed limits on table partition sizes, which allows for massive tables and consolidation of data, but this also amplifies the need for a robust ER configuration to handle the corresponding data volume. Monitoring tools, such as the enhanced InformixHQ graphical tool introduced in Informix 15, become essential for managing and tuning the ER environment at this scale.

Pros and Cons

The primary advantage of ER is its unparalleled flexibility. It supports complex, multi-node topologies, allows the target server to be fully active for both read and write operations (in an update-anywhere setup), and enables granular control over the data being replicated. This makes it ideal for workload distribution and creating active reporting servers. The asynchronous nature is also a pro for the primary server's performance, as it does not wait for remote commits.

However, this same asynchronous nature is its main drawback from a consistency standpoint. There is an inherent latency between the source and target, meaning the target server is always slightly behind the primary. This "eventual consistency" model may not be suitable for applications requiring

zero data loss failover. Furthermore, ER is more complex to set up and manage than HDR. Update-anywhere configurations introduce the possibility of data conflicts, which, while manageable through built-in resolution rules, add another layer of administrative overhead.

Data Consistency Guarantees and Downtime

Enterprise Replication guarantees transactional consistency. It replicates fully committed transactions in the order they were committed on the source, ensuring that the target database remains a transactionally consistent, albeit slightly delayed, copy of the source. It does not provide the zero-data-loss guarantee of synchronous HDR. In the event of a primary server failure, any transactions residing in the logical logs or send queue that have not yet been transmitted to the target will be lost.

The initial setup of ER on a 15TB database requires a separate data synchronization process, which implies a period of planned downtime or careful coordination to establish a baseline. The process typically involves taking the tables offline for replication, performing an initial data copy using utilities like the High-Performance Loader (HPL), and then starting ER to apply subsequent changes. Minimizing this downtime is a major planning consideration, as discussed in the later section on initial synchronization.

3. High-Availability Data Replication (HDR)

High-Availability Data Replication (HDR) is Informix's primary solution for providing high availability and disaster recovery through the creation of a hot standby server. Unlike the logical replication of ER, HDR operates at a physical level, shipping logical log records from a primary server to a secondary server, which remains in a perpetual state of recovery, applying the logs as they arrive. This ensures the secondary is a near-identical, block-for-block copy of the primary.

Architecture and Synchronization Modes

The HDR architecture consists of a single primary server and a single secondary server. The primary server performs all read-write operations and sends a continuous stream of its logical log pages to the HDR secondary. The secondary server cannot be used for read or write operations while it is acting as a standby; its sole purpose is to apply the incoming log stream and be ready to take over if the primary fails.

HDR offers two distinct synchronization modes that fundamentally define the trade-off between performance and data consistency:

Synchronous Mode: In this mode, a transaction commit on the primary server is not considered complete until the corresponding logical log records have been successfully sent to the secondary server and written to the secondary's log buffer. This creates a performance dependency, as the primary must wait for acknowledgment from the secondary. The significant benefit of this mode is that it guarantees zero data loss in the event of a primary server failure. Because the secondary is fully synchronized, it can be promoted to the primary server instantly with no loss of committed transactions. Checkpoints are also synchronized between the primary and secondary, ensuring physical consistency.

Asynchronous Mode: In this mode, the primary server commits transactions locally without waiting for acknowledgment from the secondary server. The log records are buffered and sent to the secondary as network bandwidth allows. This mode dramatically reduces the performance impact on the primary server, as it is decoupled from the secondary's status and network latency. However, this performance gain comes at the cost of potential data loss. If the primary server fails, any transactions that were committed locally but not yet transmitted to the secondary will be lost. The amount of

potential data loss is a function of the network latency and the `DRINTERVAL` setting, which controls checkpoint synchronization.

Failover Capabilities

Failover in an HDR pair is a core feature. With synchronous replication, the secondary server is a “hot” standby, ready for immediate promotion with no data loss. In cluster environments, this failover can be automated using a connection manager or other cluster software. Because the secondary is an exact physical copy, the transition is seamless for applications once they are redirected to the new primary. In asynchronous mode, failover is also possible, but it carries the risk of data loss as previously described. When the failed primary is brought back online, it can be resynchronized and assume the role of the new secondary, allowing the system to return to a protected state.

Pros and Cons

The main advantages of HDR are its simplicity and high level of data consistency. It is relatively easy to configure and manage compared to ER. In synchronous mode, it provides the strongest data protection guarantee available in Informix, making it the gold standard for mission-critical applications where data loss is unacceptable. The failover process is well-defined and reliable.

The cons are significant, however. The secondary server is a passive standby and cannot be used for any other purpose, such as running reports, which can be seen as an inefficient use of hardware resources. The one-to-one architecture (one primary to one secondary) is limiting compared to the flexible topologies of ER. The most critical drawback is the performance impact of synchronous mode on the primary server, which may be prohibitive for a high-throughput 15TB OLTP system. Asynchronous mode mitigates the performance impact but sacrifices the zero-data-loss guarantee.

4. Remote Standalone Secondary (RSS)

Remote Standalone Secondary (RSS) servers were introduced as a more flexible and scalable alternative to HDR for disaster recovery scenarios. While based on the same underlying log-shipping technology as HDR, RSS servers are designed specifically for asynchronous replication to one or more remote nodes, often in different geographical locations.

Features and Use Cases

The defining feature of RSS is its purely asynchronous communication framework. Unlike HDR, RSS does not synchronize transaction commits or checkpoints between the primary and secondary servers. The primary server sends its logical logs to the RSS nodes, but it does not wait for any acknowledgment. This design ensures that the performance of the primary server is almost entirely unaffected by the status of the RSS servers.

An Informix instance can support multiple RSS nodes, allowing for a one-to-many replication topology for disaster recovery purposes. This makes RSS ideal for maintaining off-site copies of the database in different data centers. RSS nodes also support a “delayed apply” feature, where the application of incoming logs can be intentionally postponed. This can be a valuable defense against logical data corruption on the primary; if an erroneous transaction corrupts the primary data, the delay provides a time window to stop the RSS apply process before the corruption is replicated, allowing for recovery from the secondary.

Comparison with HDR

When compared to HDR, RSS prioritizes primary server performance and disaster recovery flexibility over the immediate failover and strict consistency offered by HDR. RSS servers cannot operate in

synchronous mode and do not support automated failover via the `DRAUTO` parameter. In the event of a primary failure, an RSS server must be manually promoted, and even then, it can only be promoted to an HDR secondary, not directly to a primary server. A separate step is then required to transition it to a primary.

In a high-availability cluster containing both an HDR secondary and one or more RSS secondaries, Informix will always prioritize keeping the HDR secondary up to date. If the HDR secondary goes offline, the primary will continue sending logs to the RSS nodes. However, when the HDR node comes back online, the system will prioritize catching it up, even if it means delaying log shipments to the RSS nodes. This hierarchy exists because the HDR secondary is the designated failover target, and its log position must always be ahead of any RSS node to ensure a consistent cluster state after a failover event. This makes RSS a complementary technology to HDR, rather than a direct replacement, best used for geographically dispersed disaster recovery copies where some data loss is acceptable and immediate failover is not the primary requirement.

5. Change Data Capture (CDC) and Streaming Technologies

Beyond full database replication, Informix provides mechanisms for capturing and streaming individual data changes as they occur. These Change Data Capture (CDC) technologies are not designed to create a fully replicated database instance but are instead used to feed real-time data changes to other applications, data warehouses, or message queues. The two primary native technologies for this are Smart Triggers (also known as Push Data Triggers) and Asynchronous Post Commit Triggers.

Smart Triggers (Push Data Triggers)

Smart Triggers are an advanced CDC mechanism that allows an external client application to subscribe to changes in a dataset. The client registers its interest by providing a standard SQL `SELECT` statement with a `WHERE` clause. Internally, Informix uses components of the Enterprise Replication engine, specifically the log snoopers, to asynchronously monitor the logical logs for any committed transactions that affect rows matching the client's criteria.

When a matching change is detected, the Informix server pushes the data to the registered client. This data is delivered in a developer-friendly JSON format, which includes metadata about the transaction (e.g., operation type, commit time, transaction ID) and the row data itself (pre- and post-images of the columns). This "push" model is highly scalable because it eliminates the need for the client application to constantly poll the database, thereby reducing network traffic and database overhead. The asynchronous log-based architecture ensures that there is minimal performance impact on the primary OLTP applications that are modifying the data.

Configuration requires granting `replication` privileges to the user, ensuring a storage pool is available for Enterprise Replication to spool notices, and then using a series of `task()` functions to open a session, register for specific table events with filters, and retrieve the resulting JSON data. Sessions can be made "detachable," allowing a client to disconnect and later reconnect to resume the data stream without loss.

Asynchronous Post Commit Triggers

Asynchronous Post Commit Triggers are another mechanism for capturing data changes. While less detail is available from the provided research, they function conceptually as triggers that execute after a transaction has fully committed. Unlike standard triggers that fire as part of the transaction itself, these triggers run in a separate context, reducing their impact on the original transaction's perform-

ance. They can be used to invoke procedures or insert records into a queueing table, which can then be processed by an external application. This method provides a way to react to committed changes without the client subscription model of Smart Triggers, but it typically requires more custom development to manage the queueing and processing of the captured events. Both technologies serve the purpose of real-time data streaming but are not standalone solutions for full 15TB database replication.

6. Third-Party Solutions

While Informix's native replication tools are powerful, third-party software vendors offer solutions that address specific challenges, particularly in heterogeneous environments or where ease of use is a high priority. These tools often provide a unified interface for managing replication across different database platforms and may offer more advanced monitoring and recovery features.

Stelo and Other Commercial Tools

Stelo (formerly StarQuest Data Replicator or SQDR) is a prominent third-party replication solution that supports IBM Informix as both a source and a destination. Stelo is a Change Data Capture (CDC) based tool that typically reads from the source database's transaction logs to capture changes in real-time. It is designed to be vendor-agnostic, facilitating data movement between a wide array of relational and non-relational databases, cloud platforms like AWS and Google Cloud, and streaming technologies like Apache Kafka.

A key advantage of Stelo is its focus on minimizing the impact on production systems. It is engineered for a very low CPU footprint (often cited as less than 1%) and generally does not require any software to be installed on the source or destination database servers themselves. The replication process is managed from a central server. This approach is particularly beneficial for large-scale, high-volume environments like a 15TB Informix database, where any additional overhead on the primary server is a concern.

Stelo promotes ease of use through a graphical user interface and configuration wizards, which can simplify the setup and management of replication streams compared to the command-line heavy configuration of native tools like ER. It also includes features for automatic recovery from network interruptions, aiming to resume replication without data loss or the need for a full re-synchronization. Some users have reported that solutions like Stelo offer better visibility and more reliable performance compared to older versions of Informix ER, which could sometimes stall or have support challenges. For a large-scale Informix-to-Informix replication, Stelo presents a viable, albeit commercial, alternative that promises high reliability, low overhead, and simplified management.

7. Hybrid Approaches

Given the distinct strengths and weaknesses of each replication technology, a hybrid approach that combines two or more methods can often provide the most comprehensive and resilient solution. This is particularly true for a demanding 15TB environment that requires both high availability and the ability to offload workloads.

Combining HDR and Enterprise Replication (HDR + ER)

A powerful hybrid strategy involves using High-Availability Data Replication (HDR) and Enterprise Replication (ER) simultaneously. In this configuration, HDR is used to create a hot standby for disaster recovery and high availability. The HDR secondary remains a passive, physically identical copy of the

primary, ready for immediate, zero-data-loss failover (if using synchronous mode). This setup provides the highest level of data protection.

At the same time, Enterprise Replication is configured to replicate a subset (or all) of the data from the primary server to a third Informix server. This third server, the ER target, can be used as an active reporting or development database. It receives a continuous stream of logical updates from the primary, allowing it to serve read-intensive queries without impacting the performance of the primary OLTP server. This ER target provides the workload distribution benefit that HDR alone cannot.

This HDR + ER combination effectively separates the concerns of disaster recovery and workload balancing. HDR provides the mission-critical failover capability, while ER provides a usable, up-to-date copy of the data for other business purposes. This represents the best of both worlds, leveraging the physical robustness of HDR and the logical flexibility of ER. For a 15TB system, this approach maximizes the return on investment in hardware by having both a hot standby and an active reporting server, while ensuring the primary OLTP system is fully protected. The complexity of managing two separate replication streams is the primary trade-off.

8. Performance Considerations for 15TB

Successfully replicating a 15TB database in real-time is less about the replication technology itself and more about the underlying infrastructure's ability to support the massive data flow. Several key areas require meticulous planning and significant investment.

Network Bandwidth: This is arguably the most critical factor. The replication process, whether via HDR, RSS, or ER, involves shipping logical log data across the network. A 15TB database under heavy transactional load can generate many gigabytes of log data per hour. The network link between the primary and secondary servers must have sufficient bandwidth and low latency to handle this volume continuously. An undersized network will become a bottleneck, causing replication lag to grow, queues to fill, and potentially leading to a complete halt in replication. For synchronous HDR, high network latency will directly increase transaction commit times on the primary server.

Logical Log Sizing: The logical logs are the heart of Informix replication. For a 15TB database, the logical log space must be sized generously. A high transaction rate will cause the logs to cycle very quickly. If the replication process (ER's capture or HDR's shipping) cannot keep up with the rate of log generation, the primary server could be forced to wait for a free log, effectively stalling the entire database. It is essential to have a large number of logs and to size them appropriately to provide a sufficient buffer for the replication agents to do their work, especially during peak load or temporary network slowdowns.

Disk Space Planning: Replication requires significant additional disk space beyond the 15TB database itself. For Enterprise Replication, large send and receive queue spooling areas are needed to buffer in-flight transactions. These queues must be large enough to hold hours or even days of transactional data to survive extended network outages without losing replication data. HDR does not require spooling queues in the same way, but the secondary server must have identical disk space and layout as the primary. All replication methods also have overhead for metadata tables and internal structures.

Transaction Throughput: The design of the application's transactions has a major impact on replication performance. A small number of very large transactions (e.g., batch updates affecting millions of rows) can be more challenging for replication systems than a high volume of small, fast transactions. Large transactions generate massive logical log records that must be packaged, sent, and applied as a single unit, which can create temporary bottlenecks in the replication pipeline.

9. Initial Data Synchronization Strategies

The greatest challenge in setting up replication for a 15TB database is the initial data load. Moving this volume of data to seed the secondary server must be done as quickly as possible and with minimal disruption to the primary production system.

The most effective and recommended tool for this task is the **Informix High-Performance Loader (HPL)**. Unlike simpler utilities like `LOAD` or `dbload`, which are often limited by a 2GB file size cap and single-threaded performance, HPL is specifically designed for loading terabyte-scale data volumes. HPL's key advantage is its ability to use parallel processing. It can leverage multiple CPUs and fragment a table load across multiple threads, dramatically increasing throughput. For a 15TB database, a parallel load strategy is not optional; it is a requirement. HPL can achieve load rates of approximately 2GB per CPU per hour.

HPL's "Express Mode" is particularly well-suited for a large initial load. In this mode, HPL bypasses triggers, constraints, and indexes, performing a "light append" that writes data directly into new extents. This is significantly faster than a standard SQL insert. The methodology for a minimal downtime synchronization would be:

1. Perform a level-0 (full system) backup of the primary database.
2. Restore this backup onto the secondary server hardware. This brings the secondary to a consistent point-in-time state.
3. While the restore is happening, configure ER or HDR on the primary server to begin capturing transactions that occur after the backup was taken. These changes will be spooled.
4. Once the restore on the secondary is complete, bring the secondary server online in replication mode. It will then connect to the primary and begin receiving and applying the stream of spooled transactions, eventually catching up to the primary server.

For ER, an alternative involves using HPL to unload data from the primary and load it into the secondary. During this process, replication on the affected tables would be suspended using commands like `BEGIN WORK WITHOUT REPLICATION` for setup tasks. Once the initial data is in place and verified, ER can be started to bring the tables into sync. Regardless of the exact method, HPL is the critical component for moving the initial 15TB data set in a reasonable timeframe.

10. Best Practices

Implementing and maintaining a robust replication environment for a 15TB database requires adherence to a strict set of best practices covering configuration, monitoring, and recovery planning.

Configuration Recommendations:

- * **Isolate Replication Network:** Use a dedicated, private, high-speed network for all replication traffic. This prevents contention with application and user traffic and provides a predictable performance baseline.
- * **Optimize Disk I/O:** Place logical logs and ER spooling queues on the fastest available disks (ideally SSDs) with their own dedicated I/O controllers. This minimizes I/O contention, which is a common source of replication latency.
- * **Disable Logging on Temp Tables:** For all secondary server types (HDR, RSS), logical logging on temporary tables should be disabled by setting the `TEMPTAB_NOLOG` configuration parameter to 1. This prevents unnecessary log traffic from being generated and replicated.
- * **Tune TCP Buffers:** In high-latency network environments, tuning network parameters like

`SMX_NUMPIPES` can improve throughput by managing the number of TCP buffers in flight between the servers.

Monitoring:

- * **Proactive Monitoring is Essential:** Do not wait for replication to fail. Use InformixHQ or custom SQL-based scripts to constantly monitor key metrics.
- * **Key Metrics to Watch:** Monitor the replication lag (the time difference between the last committed transaction on the primary and its application on the secondary), the depth of the ER send/receive queues, and the logical log usage percentage.
- * **Set Up Alerts:** Configure automated alerts to notify administrators when these metrics exceed pre-defined thresholds. An alert for a rapidly growing ER send queue, for example, can provide an early warning of a network problem or a performance issue on the target server.

Conflict Resolution (for ER):

- * **Choose a Resolution Strategy:** When using update-anywhere ER, a conflict resolution strategy is mandatory. The default timestamp-based resolution (newest change wins) is often sufficient, but for complex business logic, custom-written, SQL-based conflict resolution functions may be necessary.
- * **Test Conflict Scenarios:** Rigorously test all anticipated conflict scenarios to ensure the resolution rules behave as expected. Unhandled conflicts will cause ER to abort the transaction, requiring manual intervention.

Disaster Recovery:

- * **Document the Failover Process:** Have a clear, step-by-step, documented procedure for both automated and manual failover.
- * **Regularly Test Failover:** The only way to ensure the disaster recovery plan works is to test it regularly. This includes failing over to the secondary and then failing back to the original primary. These tests build confidence and uncover any issues in the procedure or configuration before a real disaster strikes.

11. Comparison Matrix

This matrix provides a side-by-side comparison of the replication strategies analyzed in this report. Ratings are based on their suitability for a 15TB Informix-to-Informix environment with real-time requirements.

Strategy	Real-Time Capability	Data Consistency	Complexity	Initial Downtime	Performance at 15TB Scale
Enterprise Replication (ER)	High (Low Latency)	Eventual	High	Medium	Good (Scalable but requires significant tuning)
HDR (Synchronous)	Highest (Real-Time)	Strongest (Zero Data Loss)	Low	Low	Poor (High impact on primary server)
HDR (Asynchronous)	High (Low Latency)	High (Potential for minor data loss)	Low	Low	Excellent (Minimal impact on primary server)
Remote Standalone Secondary (RSS)	Medium (Asynchronous)	Eventual	Medium	Low	Excellent (No impact on primary server)
CDC / Smart Triggers	Highest (Real-Time)	N/A (Partial Data)	Medium	N/A	Excellent (Low overhead)
Third-Party (e.g., Stelo)	High (Low Latency)	High (Typically transactional)	Low	Medium	Very Good (Designed for low overhead)

Rating Scale:

- * **Real-Time Capability:** How close to zero is the replication latency.
- * **Data Consistency:** The guarantee of data integrity post-failover.
- * **Complexity:** The level of effort and expertise required for setup and maintenance.
- * **Initial Downtime:** The impact of the initial data synchronization process.
- * **Performance at 15TB Scale:** Overall suitability and efficiency for a very large database.

12. Recommended Strategy

For the specific objective of replicating a 15TB Informix database to another Informix server for continuous replication, high data consistency, and minimal downtime, no single technology is a perfect fit. Therefore, the recommended approach is a **hybrid strategy combining High-Availability Data Replication (HDR) in asynchronous mode with Enterprise Replication (ER)**.

This dual-pronged strategy is recommended for the following reasons:

1. **Robust Disaster Recovery with HDR:** HDR in asynchronous mode provides an excellent balance of performance and data protection. It creates a physically identical standby server with minimal performance overhead on the primary 15TB OLTP system. While there is a minor risk of data loss compared to synchronous mode, this risk is often acceptable when weighed against the significant performance penalty of synchronous replication at this scale. This HDR link fulfills the core requirement for a reliable, high-availability failover solution.
2. **Workload Offloading with ER:** Concurrently, Enterprise Replication should be configured to replicate data from the primary to a separate, third server. This ER target server becomes an active, readable copy of the database. It can be used to offload resource-intensive reporting, business intelligence queries, and other read-only workloads. This frees up critical resources on the primary OLTP server, allowing it to dedicate its power to processing transactions. This addresses the major limitation of HDR, where the secondary server is passive.
3. **Minimized Downtime with HPL:** The initial seeding of the 15TB of data for both the HDR and ER secondaries should be accomplished using a backup-and-restore method for the HDR node, and the **High-Performance Loader (HPL)** for the ER node. This is the only feasible method for moving a dataset of this magnitude within a reasonable time window, ensuring the initial setup downtime is minimized.

This hybrid approach provides a comprehensive solution that delivers both mission-critical disaster recovery through HDR and operational flexibility through ER. It is a complex architecture that requires careful planning and skilled administration but offers the resilience and performance necessary to support a 15TB Informix environment effectively.

References

1. [Enterprise Replication Guide - cursor-distribution.de](https://www.cursor-distribution.de/aktuell.12.10.xC1/documentation/ids_erp_bookmap.pdf) (https://www.cursor-distribution.de/aktuell.12.10.xC1/documentation/ids_erp_bookmap.pdf)
2. [IBM Informix 15: Unparalleled Scalability for the Modern, Data-Driven World - IBM](https://www.ibm.com/new/announcements/informix-15-unparalleled-scalability-modern-data-driven-world) (https://www.ibm.com/new/announcements/informix-15-unparalleled-scalability-modern-data-driven-world)
3. [Enterprise Replication - IBM](https://www.ibm.com/docs/en/informix-servers/14.10?topic=replication-enterprise) (https://www.ibm.com/docs/en/informix-servers/14.10?topic=replication-enterprise)
4. [Informix Enterprise Replication: A Technical Overview - IBM](https://www.ibm.com/docs/en/informix-servers/12.10.0?topic=replication-informix-enterprise-technical-overview) (https://www.ibm.com/docs/en/informix-servers/12.10.0?topic=replication-informix-enterprise-technical-overview)
5. [Comparison of RS secondary servers and HDR secondary servers - IBM](https://www.ibm.com/docs/en/informix-servers/12.10?topic=servers-comparison-rs-secondary-hdr-secondary) (https://www.ibm.com/docs/en/informix-servers/12.10?topic=servers-comparison-rs-secondary-hdr-secondary)
6. [RSS with delay apply is dog slow - comp.databases.informix.narkive.com](https://comp.databases.informix.narkive.com/BTljCl37/rss-with-delay-apply-is-dog-slow) (https://comp.databases.informix.narkive.com/BTljCl37/rss-with-delay-apply-is-dog-slow)
7. [Always On - High Availability Best Practices for Informix - slideshare.net](https://www.slideshare.net/slideshow/always-on-high-availability-best-practices-for-informix/55285704) (https://www.slideshare.net/slideshow/always-on-high-availability-best-practices-for-informix/55285704)
8. [All about Informix Data Replication - beinghealthy602843940.wordpress.com](https://beinghealthy602843940.wordpress.com/2019/09/01/all-about-informix-data-replication/) (https://beinghealthy602843940.wordpress.com/2019/09/01/all-about-informix-data-replication/)
9. [RSS - docs.deistercloud.com](https://docs.deistercloud.com/content/Databases.30/IBM%20Informix.2/Replication/RSS.xml?embedded=true) (https://docs.deistercloud.com/content/Databases.30/IBM%20Informix.2/Replication/RSS.xml?embedded=true)

10. [Informix TechTalk: Capturing Data Changes with Art Kagel - The International Informix Users Group](https://www.iiug.org/en/2021/02/25/informixtechtalk-20210225-capturingdatachanges-artkagel/) (<https://www.iiug.org/en/2021/02/25/informixtechtalk-20210225-capturingdatachanges-artkagel/>)
11. [InformixTechTalk-20210225-CapturingDataChanges-ArtKagel.pdf - The International Informix Users Group](https://www.iiug.org/wp-content/uploads/2021/02/InformixTechTalk-20210225-CapturingDataChanges-ArtKagel.pdf) (<https://www.iiug.org/wp-content/uploads/2021/02/InformixTechTalk-20210225-CapturingDataChanges-ArtKagel.pdf>)
12. [Change Data Capture API - IBM](https://www.ibm.com/docs/en/informix-servers/12.10.0?topic=api-change-data-capture) (<https://www.ibm.com/docs/en/informix-servers/12.10.0?topic=api-change-data-capture>)
13. [Informix TechTalk: Capturing Data Changes - YouTube](https://www.youtube.com/watch?v=ion1JtLzNQI) (<https://www.youtube.com/watch?v=ion1JtLzNQI>)
14. [Change Data Capture Overview - cursor-distribution.de](https://www.cursor-distribution.de/aktuell.11.50.xC9/documentation/ids_cdc_bookmap.pdf) (https://www.cursor-distribution.de/aktuell.11.50.xC9/documentation/ids_cdc_bookmap.pdf)
15. [Informix Webcast: Loading Data into Informix - FAQ - advancedatatools.com](https://advancedatatools.com/Downloads/Informix-Webcast_Load_FAQ.htm) (https://advancedatatools.com/Downloads/Informix-Webcast_Load_FAQ.htm)
16. [How to unload and load millions of records very fast into Informix. - YouTube](https://www.youtube.com/watch?v=FFy-YTeds-g) (<https://www.youtube.com/watch?v=FFy-YTeds-g>)
17. [Using the High-Performance Loader to load and unload data from a OneDB database - Actian Academy](https://academy.actian.com/using-the-high-performance-loader-to-load-and-unload-data-from-a-onedb-database) (<https://academy.actian.com/using-the-high-performance-loader-to-load-and-unload-data-from-a-onedb-database>)
18. [informix load from file with header - Stack Overflow](https://stackoverflow.com/questions/28478657/informix-load-from-file-with-header) (<https://stackoverflow.com/questions/28478657/informix-load-from-file-with-header>)
19. [IBM Informix High-Performance Loader User's Guide - IBM](https://publibfp.dhe.ibm.com/epubs/pdf/c2394332.pdf) (<https://publibfp.dhe.ibm.com/epubs/pdf/c2394332.pdf>)
20. [Stelo software as a replacement for Enterprise Replication? - IBM Community](https://community.ibm.com/community/user/discussion/stelo-software-as-a-replacement-for-enterprise-replication) (<https://community.ibm.com/community/user/discussion/stelo-software-as-a-replacement-for-enterprise-replication>)
21. [Informix Real-Time Transaction Replication, Migration & Streaming - Stelo](https://www.stelodata.com/technologies/informix) (<https://www.stelodata.com/technologies/informix>)
22. [Stelo releases V6.51; Web-based GUI delivers simplified, scalable data streaming for legacy-to-AI pipelines - PR Newswire](https://www.prnewswire.com/news-releases/stelo-releases-v6-51-web-based-gui-delivers-simplified-scalable-data-streaming-for-legacy-to-ai-pipelines-302453469.html) (<https://www.prnewswire.com/news-releases/stelo-releases-v6-51-web-based-gui-delivers-simplified-scalable-data-streaming-for-legacy-to-ai-pipelines-302453469.html>)
23. [Stelo Data Replicator v 6.70 Readme - Stelo](https://docs.stelodata.com/Products/sqdr/readme) (<https://docs.stelodata.com/Products/sqdr/readme>)
24. [Stelo enables mirroring and streaming data replication into Google Big Query and Databricks from Db2, Oracle, and SQL Server in latest release - PR Newswire](https://www.prnewswire.com/news-releases/stelo-enables-mirroring-and-streaming-data-replication-into-google-big-query-and-databricks-from-db2-oracle-and-sql-server-in-latest-release-301998722.html) (<https://www.prnewswire.com/news-releases/stelo-enables-mirroring-and-streaming-data-replication-into-google-big-query-and-databricks-from-db2-oracle-and-sql-server-in-latest-release-301998722.html>)